

Project Report CS 553 Fall 2015

Overview

In this project we have built a search engine. This search engine could be used by a user of the internet to search for web pages he/she is looking for. The user can provide words associated with the web page he/she is looking for as the input, and the search engine looks for the words entered by the user in the corpus (collection of the parsed data of web pages that the search engine has) and returns the most appropriate results (URLs of web-pages which contain the word(s) provided as input) ranked using the Ranking Algorithm explained in the report. The project is divided in three phases which are explained in the project later.

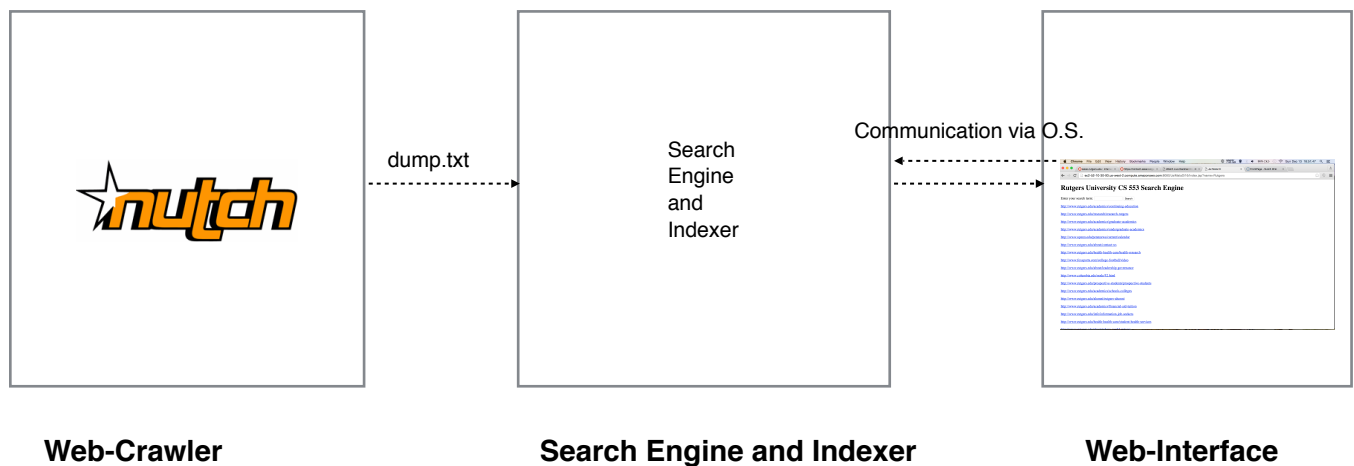
This report is organized as follows -

- Overview of Phases
 1. Phase 1 - Crawling using Nutch
 - 1.1 Crawl Database Analysis and Discussion
 2. Phase 2 - Search Engine and Indexer
 - 2.1 Development Environment of the Search Engine and Indexer
 - 2.2 File Processing
 - 2.3 Token Generation
 - 2.4 Stemming
 - 2.5 Index Generation
 - 2.6 Scoring
 3. Phase 3 - Web Interface
 4. Flow and Algorithm
 5. General Observations
- Summary of Contributions

Overview of Phases

This project was divided in 3 phases, each of which maps to a specific module of the search engine -

1. Phase 1 - [nutch](#) a web-crawler was used to crawl the internet. It accepted seed lists consisting of educational websites and sports websites. The output of this phase was dump.txt file which had 76.2 MB size and contained more than 1.6 million lines containing all the fetched web pages, urls, and metadata.
2. Phase 2 - This phase accepts as input the dump.txt file generated as output to phase 1. It then parses the dump.txt file extracts all the parsed data of the fetched web-pages and indexes all the web-pages. The index is an inverted-index (which means that it is a key-value index where the key is the word and the value is the document number which contains the word). In this phase a Ranking Algorithm was implemented to rank the outputs. The input to the Indexer and Search Engine is the search term (consisting of words being looked for) and output is the ranked list of URLs of web-pages containing that word.
3. Phase 3 - Web-Interfacing. A Web-Interface ([website](#)) which allows the users to enter the search term, which is accepted by the Indexer and Search Engine (described in Phase 2) as the input, and then subsequently displays the output from the Indexer and Search Engine on the web-interface. The interface exposes the search engine to a user of the internet, making it possible for a common user to use it.



Phase 1 - Crawling using Nutch

In Phase 1 we used [Nutch 1.0](#) crawler to crawl the web. Apache Nutch is a highly extensible and scalable open source web crawler software. The *input* for this phase were 2 seed lists. Seed lists consisted of the urls, these urls act as starting points for crawling. We used the [nutch tutorial](#) for this purpose.

The internet can be compared to a giant graph where each of the computers connected to it, is like a node of the graph. Since nutch by default crawls in Breadth-First-Search manner, therefore it was natural to provide a seed list. This is because [BFS](#) is a graph traversal algorithm which accepts the source node as input for it to begin traversing the graph. BFS begins traversal from the source node provided to it as the input. There is a second reason to provide the seed list - if specific websites need to be crawled then providing their exact urls is useful. Crawling the web is like traversing the giant-graph called internet, its just that, while crawling whatever the crawler hits upon (e.g. a web page) it fetches that, so that it could be used for other purpose (e.g. input for a search engine) later on.

We provide here the list of educational websites and sports websites which constituted the seed lists

<http://www.harvard.edu/>
<http://www.princeton.edu/>
<http://www.yale.edu/>
<http://www.columbia.edu/>
<http://www.caltech.edu/>
<http://web.mit.edu/>
<http://www.rutgers.edu/>
<http://www.stanford.edu/>
<http://www.uchicago.edu/>
<http://www.upenn.edu/>
<http://www.duke.edu/>
<http://www.dartmouth.edu/>
<http://www.northwestern.edu/>
<http://www.jhu.edu/>
<http://wustl.edu/>
<http://www.brown.edu/>
<http://www.rice.edu/>
<http://www.vanderbilt.edu/>
<http://nd.edu/>
<http://www.emory.edu/home/>
<http://www.asu.edu/>
<http://www.washington.edu/>
<http://www.foxsports.com/>
<http://www.nbcsports.com/>
<http://www.bbc.com/sport/0/>
<http://www.realmadrid.com/en>
<http://www.wimbledon.com/>
<http://www.usopen.org/index.html>
<http://www.sportstats.com>
<http://bleacherreport.com/nfl>
<http://www.nfl.com/>
<http://www.nblcanada.com/>
<http://www.ussoccer.com/>
<http://www.fifa.com>
<http://www.uefa.com>
<http://www.usbasket.com/>
<http://princetonyouthhockey.org/>
<https://www.nsaa.org/>
<http://www.isaa.org/>

As mentioned in the Nutch Tutorial the crawler - Nutch was used to fetch 2 times after going through the following steps

1. Seeding
2. Fetching
3. Parsing
4. Updating the DB
5. Fetching 2nd time
6. Parsing 2nd time
7. Updating DB 2nd time
8. Inverting all links

All the above steps are explained in detail in the nutch tutorial. We present here the significant observations. We show the commands and discuss and analyze the output of the commands below.

1.1 Crawl Database Analysis and Discussion

Universities Websites- It was observed that 3733 links were there. We have given below the command and the output.

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10$ bin/nutch readdb crawl/crawldb/ -stats
CrawlDb statistics start: crawl/crawldb/
Statistics for CrawlDb: crawl/crawldb/
TOTAL urls: 3733
retry 0: 3732
retry 1: 1
min score: 0.0
avg score: 0.012907581
max score: 1.552
status 1 (db_unfetched): 3147
status 2 (db_fetched): 497
status 3 (db_gone): 6
status 4 (db_redir_temp): 13
status 5 (db_redir_perm): 70
CrawlDb statistics: done
```

Sports Websites - It was observed that 2907 links were there. We have given below the command and the output.

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10$ bin/nutch readdb crawl/crawldb/ -stats
CrawlDb statistics start: crawl/crawldb/
Statistics for CrawlDb: crawl/crawldb/
TOTAL urls: 2907
retry 0: 2902
retry 1: 5
min score: 0.0
avg score: 0.014617475
max score: 1.604
status 1 (db_unfetched): 2316
status 2 (db_fetched): 533
status 3 (db_gone): 2
status 4 (db_redir_temp): 48
status 5 (db_redir_perm): 8
CrawlDb statistics: done
```

Understanding the Output

The ReadDb command is used read the crawl dump.

EDUCATION WEBSITES

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10$ bin/nutch readdb crawl/crawldb/ -dump joutput
CrawlDb dump: starting
CrawlDb db: crawl/crawldb/
CrawlDb dump: done
```

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10$ bin/nutch invertlinks crawl/linkdb -dir crawl/segments
```

```
LinkDb: starting at 2015-10-12 04:45:28
LinkDb: linkdb: crawl/linkdb
LinkDb: URL normalize: true
LinkDb: URL filter: true
LinkDb: external links will be ignored.
LinkDb: adding segment: file:/home/jaimatadi/Documents/JaiMataDi/apache-nutch-1.10/crawl/segments/20151012040815
LinkDb: adding segment: file:/home/jaimatadi/Documents/JaiMataDi/apache-nutch-1.10/crawl/segments/20151012040927
LinkDb: finished at 2015-10-12 04:45:32, elapsed: 00:00:04
```

SPORTS WEBSITES

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10$ bin/nutch readdb crawl/crawldb/ -dump
joutput
CrawlDb dump: starting
CrawlDb db: crawl/crawldb/
CrawlDb dump: done
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10$ bin/nutch readlinkdb crawl/linkdb/ -dump
jlinkdb_output
LinkDb dump: starting at 2015-10-12 10:13:20
LinkDb dump: db: crawl/linkdb/
LinkDb dump: finished at 2015-10-12 10:13:22, elapsed: 00:00:01
```

Analysis -

As a result of the **ReadDb** command a file called **part-00000** is created. For the education sites the size of the file was 1.1 MB for the Sports website the size was 902.6 KB.

When we crawl the output is in the **crawl/crawldb** directory. However, this output is not human readable. It is inside a directory called part-0000 which has data and index files.

But due to the ReadDb command we see human understandable output. This is easy for text processing. The file contains details about url, status, metadata and so on.

Excerpt from the part-0000 file -

```
http://nd.edu/ Version: 7
Status: 5 (db_redir_perm)
Fetch time: Wed Nov 11 03:08:40 EST 2015
Modified time: Wed Dec 31 19:00:00 EST 1969
Retries since fetch: 0
Retry interval: 2592000 seconds (30 days)
Score: 1.0
```

Signature: null
Metadata:
Content-Type=text/html
pst=moved(12), lastModified=0: <http://www.nd.edu/>
rs=38

<http://web.mit.edu/> Version: 7
Status: 2 (db_fetched)
Fetch time: Wed Nov 11 03:08:40 EST 2015
Modified time: Wed Dec 31 19:00:00 EST 1969
Retries since fetch: 0
Retry interval: 2592000 seconds (30 days)
Score: 1.0734382
Signature: 6359cab4384a5706c108773b391e09b
Metadata:
Content-Type=text/html
pst=success(1), lastModified=0
rs=45

<http://web.mit.edu/aboutmit> Version: 7
Status: 5 (db_redir_perm)
Fetch time: Wed Nov 11 03:09:49 EST 2015
Modified time: Wed Dec 31 19:00:00 EST 1969
Retries since fetch: 0
Retry interval: 2592000 seconds (30 days)
Score: 0.039081886
Signature: null
Metadata:
Content-Type=text/html
pst=moved(12), lastModified=0: <http://web.mit.edu/aboutmit/>
rs=101

<http://web.mit.edu/aboutmit/> Version: 7
Status: 1 (db_unfetched)
Fetch time: Mon Oct 12 04:29:19 EDT 2015
Modified time: Wed Dec 31 19:00:00 EST 1969
Retries since fetch: 0
Retry interval: 2592000 seconds (30 days)
Score: 0.01135637
Signature: null
Metadata:

As a result of the **invertlink** command the links were inverted and **linkDb** was created. Now using the **readlinkdb**. Using the **readlinkdb** command an output file **part-00000**. This file contains information about inlinks. For education websites the size of this file is 2.1 MB and for Sports websites it was 2.8 MB.

This data is obtained from the **crawl/linkdb** directory which has current directory with an index and data file. Both are in non-human readable format. The commands convert this data to human understandable format which is easy for text processing.

An excerpt of the file for Sports Websites is.

<http://bleacherreport.com/articles/feed> Inlinks:
fromUri: <http://bleacherreport.com/nfl> anchor:

<http://www.fiba.com/> Inlinks:
fromUri: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/CRO/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUri: <http://www.fiba.com/media-centre/home> anchor: FIBA.com
fromUri: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/BRA/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUri: <http://www.fiba.com/oceaniaawomen/2015> anchor: FIBA.com
fromUri: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/WNBA/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUri: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/IRI/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUri: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/LEBA/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUri: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/BIH/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUri: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/AUT/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUri: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/WHUN/ot/1/leagueProfile.html> anchor: FIBA Logo

fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/BLR/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/EL/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/GER/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/WCZE/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/rankinggirls> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/INA/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/BEL/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/BALKAN/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/ARG/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/ISR2/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/WABA/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/ASEAN/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/world/u19women/2015> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/ISL/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/TPE/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/EST/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/DEN/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/NBA/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/world/u19/2015> anchor: FIBA.com
fromUrl: <http://www.fiba.com/rankingcombined> anchor: FIBA.com
fromUrl: <http://www.fiba.com/rankingmen> anchor: FIBA.com
fromUrl: <http://www.fiba.com/afrobasket/2015> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/MEX/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/FRA/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/world/women/2014> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/openNodeIds/21192/selNodeID/21192/allResults.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/COPCC/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/afrobasketwomen/2015> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/KAZ/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/GBR/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/PHI/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/oceania/2015> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/QAT/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/UC/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/rankingboys> anchor: FIBA.com
fromUrl: <http://www.fiba.com/3x3U18WC/2015> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/GER2/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/URU/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/ECM2/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/asiawomen/2015> anchor: FIBA.com
fromUrl: <http://www.fiba.com/3x3WT/2015/abu-dhabi> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/GEO/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/asia/2015> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/ADRIATIC/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/CAN/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/GRE/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/ELW/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/BUL/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/BAL TIC_B/ot/1/leagueProfile.html anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/SUDAMERICA/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/hall-of-fame> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/CHN/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/WCRO/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/world/u17/2014> anchor: FIBA.com
fromUrl: <http://www.fiba.com/world/u17women/2014> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/HUN/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/ASIA_CHAMPIONS/ot/1/leagueProfile.html anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/ECW/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/news> anchor: FIBA.com
fromUrl: <http://www.fiba.com/rankingwomen> anchor: FIBA.com
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/FIN/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/NBDL/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/VEN/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl:

fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/CZE/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/BALTIC/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/FRA2/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/pages/eng/fc/gameCent/p/leagueid/DOM/ot/1/leagueProfile.html> anchor: FIBA Logo
fromUrl: <http://www.fiba.com/asia/u16women/2015> anchor: FIBA.com

<http://www.fiba.com/3x3/federations-ranking> Inlinks:
fromUrl: <http://www.fiba.com/rankingwomen> anchor: Federation Ranking
fromUrl: <http://www.fiba.com/events> anchor: Federation Ranking
fromUrl: <http://www.fiba.com/asia/u16women/2015> anchor: Federation Ranking
fromUrl: <http://www.fiba.com/hall-of-fame> anchor: Federation Ranking
fromUrl: <http://www.fiba.com/world/u17women/2014> anchor: Federation Ranking
fromUrl: <http://www.fiba.com/> anchor: Federation Ranking
fromUrl: <http://www.fiba.com/world/u17/2014> anchor: Federation Ranking
fromUrl: <http://www.fiba.com/news> anchor: Federation Ranking
fromUrl: <http://www.fiba.com/3x3WT/2015/abu-dhabi> anchor: Federation Ranking
fromUrl: <http://www.fiba.com/asiawomen/2015> anchor: Federation Ranking
fromUrl: <http://www.fiba.com/asia/2015> anchor: Federation Ranking
fromUrl: <http://www.fiba.com/oceania/2015> anchor: Federation Ranking

INLINK CALCULATION

The part-0000 text files outputted due to the readlinkdb command were analyzed first visually. After visual inspection a pattern was observed and then text processing was used to find the inlinks

For the Education Sites there were 24205 unlinks

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/jlinkdb_output$ grep fromUrl: part-00000 | wc -l  
24205
```

For the Sports Sites there were 33330 unlinks

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10/jlinkdb_output$ grep fromUrl: part-00000 | wc -l  
33330
```

READSEG COMMAND and DUMPING OUTPUT of FETCHED DATA

The readseg command dumps the content of a segment. We used it to dump the segment contents in a human readable format. It was visually analyzed.

This command takes as input the crawl/segments/<TimeStamp of the segment> as input and dumps a file which can be text processed

```
bin/nutch readseg -dump crawl/segments/20151012040927/ jreadseg_output  
SegmentReader: dump segment: crawl/segments/20151012040927  
SegmentReader: done
```


For the Education Websites it read dumped a file which was 29.3 MB in size and for the Sports Websites the file was 43.3 MB. These two files were later combined to obtain the dump.txt file of size 76.2 MB, which is used as input to the Second phase.

NOTE - These files were big in size and lot of text processing was required.
These have MetaData + HTML + ParsedData of the pages fetched.

OUTLINKS CALCULATION

It was seen that the readlinkdb tool could be used to find list of inlinks. To find the out-links we need to merge the segments and read the result.

The mergesegs command was used to merge the segments.

Then the merged segments were read using readseg command

The output file was visually inspected and text processed to find the number of outlinks.

It was found that there were 28700 out-links for the education website and 36263 for the sports one.

EDUCATION WEBSITE MERGE SEGMENTS

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10$ bin/nutch mergesegs jmergesegs crawl/segments/20151012040815/ crawl/segments/20151012040927/
Merging 2 segments to jmergesegs/20151012104533
SegmentMerger: adding crawl/segments/20151012040815
SegmentMerger: adding crawl/segments/20151012040927
SegmentMerger: using segment data from: content crawl_generate crawl_fetch crawl_parse parse_data parse_text
```

SPORT SITE MERGE SEGMENTS

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10$ bin/nutch mergesegs jmergesegs crawl/segments/20151012042631/ crawl/segments/20151012042759/
Merging 2 segments to jmergesegs/20151012105857
SegmentMerger: adding crawl/segments/20151012042631
SegmentMerger: adding crawl/segments/20151012042759
SegmentMerger: using segment data from: content crawl_generate crawl_fetch crawl_parse parse_data parse_text
```

SPORTS WEBSITE OUTLINKS

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10/joutlinks$ bin/nutch readseg -dump jmergesegs/20151012105857/ joutlinks
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10/joutlinks$ grep Outlinks dump > joutlinkdetails.txt
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10/joutlinks$ awk '{s+=$0} END {print s}' joutlinkdetails.txt
36263
```

LINK DEPTHS ANALYSIS

EDUCATION WEB SITES

URLS with depth and their counts -

Depth of 8= 1

Depth of 7 9 - 1 = 8

Depth of 6 66 - 9 = 57

Depth of 5 277 - 66 = 211

Depth of 4 827 - 277 = 550

Depth of 3 1859 - 827 = 1032

Depth of 2 3523 - 1859 = 1664

Depth of 1 3787 - 3523 = 264

Total Number of URLS = 1 + 8 + 57 + 211 + 550 + 1032 + 1664 + 264 = 3787 which matches the initial statistics found using readdb.

This took as input the output of the readdb which was dumped into a text file.

Using regular expressions this was found the output and commands are below-

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/joutput$ grep -o http://[a-zA-Z.0-9+~]* / part-00000 | wc -l
3787
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/joutput$ grep -o http://[a-zA-Z.0-9+~]*[a-zA-Z0-9+~]* /
part-00000 | wc -l
3523
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/joutput$ grep -o http://[a-zA-Z.0-9+~]*[a-zA-Z0-9+~]*[a-zA-
Z0-9+~]* / part-00000 | wc -l
1859
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/joutput$ grep -o http://[a-zA-Z.0-9+~]*[a-zA-Z0-9+~]*[a-zA-
Z0-9+~]*[a-zA-Z0-9+~]* / part-00000 | wc -l
827
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/joutput$ grep -o http://[a-zA-Z.0-9+~]*[a-zA-Z0-9+~]*[a-zA-
Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]* / part-00000 | wc -l
277
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/joutput$ grep -o http://[a-zA-Z.0-9+~]*[a-zA-Z0-9+~]*[a-zA-
Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]* / part-00000 | wc -l
66
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/joutput$ grep -o http://[a-zA-Z.0-9+~]*[a-zA-Z0-9+~]*[a-zA-
Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]* / part-00000 | wc -l
9
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/joutput$ grep -o http://[a-zA-Z.0-9+~]*[a-zA-Z0-9+~]*[a-zA-
Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]* / part-00000 | wc -l
1
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/joutput$ grep -o http://[a-zA-Z.0-9+~]*[a-zA-Z0-9+~]*[a-zA-
Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]*[a-zA-Z0-9+~]* / part-00000 | wc -l
0
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/joutput$
```


10762 and
9831 the
9408 of
5952 to
4346 in
3876 for
3674 a
3462 Caltech
3183 University
2423 s
2228 Research

2005 School
1991 at
1820 on
1818 is
1719 2015
1717 The
1689 Oct
1624 Center
1475 with
1418 Campus
1334 News
1319 06
1230 Rutgers
1219 Students
1208 ly
1198 that
1180 Graduate
1126 students
1123 Penn
1117 ow
1089 Admissions
1058 Faculty
1022 Education
995 research
993 Resources
992 Flickr
992 12
988 t
963 by
959 Health
955 Academic
929 A
920 Events
905 Harvard
898 co
895 Rice
876 Alumni
873 Services
872 are
867 Undergraduate
864 Programs
855 or
854 http
837 Staff
836 from
828 About
822 New
820 as
819 Calendar
802 Sciences
755 Student
755 Our
753 Arts
748 Studies
743 Engineering
740 Office
732 more
725 our
716 Visit
701 Academics
699 Public
697 Aid
680 Search
680 Contact
658 Vanderbilt
655 an
649 you
629 Information
619 Media
618 Institute
600 Science
591 campus

587 Columbia
583 Athletics
570 Yale
570 College
567 Northwestern
567 be
563 Life
559 Directory
558 Give
533 More
533 faculty
530 will
527 their
526 Giving
523 History
522 Financial
516 about
513 Brown
511 new
507 this
502 Z
501 Maps
491 International
490 Us
488 Program
485 programs
485 Global
483 Facebook
478 Twitter
473 Social
470 Washington
462 Learning
448 world
448 edu
446 Offices
446 Libraries
445 Online
444 your
438 322
430 Library
420 Medicine
416 information
410 President
409 Community
408 Continuing
402 Leadership
400 Safety
400 content
399 Policy
399 Parents
391 community
387 has
384 all
382 Schools
374 We
363 we
363 Professional
361 Medical
359 Copyright
354 main
348 S
347 have
347 01

The example with explanation of the command used for this purpose-

The command:

```
/Downloads/apache-nutch-1.10/runtime/local/out2db$ grep -o -E 'w+' part-00000 | sort -f | uniq -c | sort -bnr
```

The explanation:

grep:

-o only prints the portion of the line that matches that pattern.
-E searches for a certain pattern.
'w+' indicates that "word" is the pattern to be matched.

sort :

-f tells sort to ignore case when comparing words.

-bnr sorts in a numeric reverse order, while ignoring whitespaces.

Uniq:

-c counts the occurrences.

SPORTS WEBSITES FREQUENCY OF WORDS -

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10/jreadseg_output$ grep -A 1 ParseText:: dump > jParsedText.txt  
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10/jreadseg_output$ vi jParsedText.txt
```

```
[3]+ Stopped            vi jParsedText.txt
```

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10/jreadseg_output$ grep -o -E 'w+' jParsedText.txt | sort -f | uniq -c | sort -bnr > jWordCount.txt  
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/JaiMataDiSports/apache-nutch-1.10/jreadseg_output$
```

OUTPUT (ONLY SOME PART IS HERE)

```
4863 the  
3417 U  
3148 of  
3053 0  
2637 1  
2427 s  
2343 and  
2141 to  
2110 Stories  
2085 About  
2039 n  
2018 Roster  
1922 2015  
1878 Matches  
1852 in  
1807 2  
1764 Stats  
1764 Latest  
1716 Tournaments  
1625 a
```

1508 Cup
1437 Coaching
1417 Women
1377 t
1359 The
1292 Soccer
1292 News
1270 5
1259 3
1231 World
1222 Staff
1218 i
1194 Tickets
1123 4
1033 for
1028 MNT
1027 League
1018 S
1009 FIFA
1006 NBLC
990 u
983 r
973 on
948 WNT
901 All
897 2014
869 15
848 UEFA
844 with
839 6
839 10
838 Home
838 f
768 20
757 Football
744 Club
743 FIBA
719 17
699 this
695 at
684 function
655 Game
654 return
643 8
629 16
625 NBSCN
622 19
620 com
620 18
616 Canada
595 ET
591 A
590 vs
586 Men
573 e
568 National
568 Islands
567 Republic
560 23
558 NFL
553 Live
539 Team
539 Basketball
533 o
532 7
531 11
525 is
518 New
505 USA
504 Search
504 NBL
501 2013

498 Real
493 9
487 by
485 Week
480 Standings
468 Store
465 Season
461 FOX
453 Liga
444 Archive
438 Teams
429 from
427 Championship
425 12
424 your
422 Members
419 Bernabéu
419 be
415 Schedule
414 14
411 Video
411 or
409 Madrid
396 Contact

HIGHEST FREQUENCY URL

We use similar technique as used for word frequency to find the highest frequency urls. The higher the frequency the more popular and the more connected. We take as input the readdb command's output's parse-0000 dump.

The example with explanation of the command used for this purpose-

The command combination:

```
grep -o -E "\<"http:/".*\>/" part-00000 | sort -f | uniq -c | sort -bnr
```

The explanation:

"\<"http:/".*\>/" will consider only the URLs as a pattern to be parsed.

EDUCATION CENTER

URL RANKING

```
jaimatadi@jaimatadi-laptop:~/Documents/JaiMataDi/apache-nutch-1.10/joutput$ grep -o -E "\<"http:/".*\>/" part-00000 | sort -f | uniq -c | sort -bnr > jLinkRank.txt
```

OUTPUT (SOME PART ONLY)

```
126 http://www.caltech.edu/content/  
79 http://www.washington.edu/students/crscat/  
76 http://www.caltech.edu/node/  
68 http://www.caltech.edu/news/  
53 http://www.vanderbilt.edu/  
46 http://www.upenn.edu/node/  
40 http://www.caltech.edu/file/  
39 http://www.upenn.edu/spotlights/  
38 http://www.columbia.edu/content/
```

36 <http://www.rutgers.edu/about/>
33 <http://www.rutgers.edu/node/>
32 http://www.caltech.edu/news/tag_ids/
32 <http://web.mit.edu/facts/>
31 <http://www.rutgers.edu/academics/>
31 <http://www.columbia.edu/node/>
31 <http://web.mit.edu/research/topic/>
28 <http://www.harvard.edu/>
25 <http://www.yale.edu/about/>
25 <http://www.washington.edu/students/reg/>
24 <http://www.vanderbilt.edu/atoz/letter/>
24 <http://www.upenn.edu/penna-z/>
22 <http://www.northwestern.edu/studyabroad/programs/europe/>
21 <http://www.vanderbilt.edu/atoz/tag/>
21 <http://www.upenn.edu/>
19 <http://www.columbia.edu/cu/opir/abstract/>
19 <http://web.mit.edu/>
18 <http://www.vanderbilt.edu/career/tools/>
18 <http://www.rice.edu/>
16 <http://www.vanderbilt.edu/career/>
16 <http://www.upenn.edu/pennnews/>
16 <http://web.mit.edu/community/topic/>
15 <http://www.washington.edu/research/SITEPARTS/documents/.or/>
15 <http://www.brown.edu/academics/engineering/about/>
15 <http://www.brown.edu/>
14 <http://www.upenn.edu/about/>
14 <http://www.caltech.edu/>
14 <http://web.mit.edu/industry/>
13 <http://www.upenn.edu/programs/>
13 <http://www.upenn.edu/pennnews/current/>
12 <http://www.yale.edu/hronline/benefits/>
12 <http://www.yale.edu/>
12 <http://www.vanderbilt.edu/financialaid/>
11 <http://www.rice.edu/unconventional/>
11 <http://www.brown.edu/academics/engineering/>
11 <http://web.mit.edu/staff/>
10 <http://www.yale.edu/gateways/>
10 <http://www.upenn.edu/pennnews/current/calendar/%40--2015-11-11/>
10 <http://www.upenn.edu/life-at-penn/>
10 <http://www.northwestern.edu/about/our-people/>
10 <http://www.harvard.edu/about-harvard/>
10 <http://www.columbia.edu/files/columbia/content/>
10 <http://www.brown.edu/about/>
10 <http://web.mit.edu/staff/business/>
9 <http://www.yale.edu/hronline/careers/>
9 <http://www.upenn.edu/sites/default/files/research-at-penn/pdf/>
9 <http://www.stanford.edu/>
9 <http://www.northwestern.edu/studyabroad/programs/americas/>
9 <http://www.brown.edu/campus-life/events/family-weekend/>
9 <http://www.brown.edu/campus-life/>
9 <http://www.brown.edu/admission/undergraduate/>
9 <http://www.brown.edu/academics/>
8 <http://www.vanderbilt.edu/work-at-vanderbilt/>
8 <http://www.upenn.edu/pennnews/experts/>
8 <http://www.upenn.edu/computing/security/advisories/>
8 <http://www.northwestern.edu/studyabroad/programs/asia/>
8 <http://www.northwestern.edu/globalhealthstudies/about/>
8 <http://www.harvard.edu/media-relations/>
8 <http://www.brown.edu/academics/engineering/about/events/2015-10/>
7 <http://www.vanderbilt.edu/financialaid/undergraduate/>
7 <http://www.upenn.edu/services/>
7 <http://www.upenn.edu/president/meet-president/>
7 <http://www.upenn.edu/computing/security/>
7 <http://www.stanford.edu/academics/>
7 <http://www.northwestern.edu/studyabroad/programs/africa/>
7 <http://www.harvard.edu/on-campus/commencement/>
7 <http://www.harvard.edu/about-harvard/harvard-glance/>
7 <http://web.mit.edu/mitpsc/whoweare/>
6 <http://www.yale.edu/medicine/>
6 <http://www.washington.edu/students/gencaat/front/>
6 <http://www.washington.edu/>
6 <http://www.upenn.edu/sustainability/sustainability-themes/>

6 <http://www.uchicago.edu/community/>
 6 <http://www.northwestern.edu/newscenter/stories/2015/10/>
 6 <http://www.northwestern.edu/globalhealthstudies/faculty/>
 6 <http://www.northwestern.edu/diversity/>
 6 <http://www.harvard.edu/visitors/>
 6 <http://www.brown.edu/gateway/>
 6 <http://www.brown.edu/academics/professional/>
 6 <http://www.brown.edu/academics/gradschool/>
 6 <http://www.brown.edu/about/administration/global-engagement/>
 6 <http://wustl.edu/academics/>
 6 <http://web.mit.edu/staff/culture/>
 5 <http://www.yale.edu/hronline/>
 5 <http://www.vanderbilt.edu/ResEd/main/>
 5 <http://www.vanderbilt.edu/nashville/>
 5 <http://www.vanderbilt.edu/isss/wp-content/uploads/>
 5 <http://www.vanderbilt.edu/international/funding/>
 5 <http://www.vanderbilt.edu/financialaid/loans-payment/>
 5 <http://www.upenn.edu/president/about-presidency/>
 5 <http://www.upenn.edu/pennnews/current/2015-10-08/latest-news/>
 5 <http://www.uchicago.edu/students/>
 5 <http://www.stanford.edu/about/>
 5 <http://www.princeton.edu/main/>
 5 <http://www.northwestern.edu/studyabroad/programs/middle-east/>
 5 <http://www.northwestern.edu/academics/>
 5 <http://www.northwestern.edu/about/>
 5 <http://www.northwestern.edu/>
 5 <http://www.harvard.edu/on-campus/visit-harvard/>
 5 <http://www.harvard.edu/on-campus/>

SPORTS CENTER RANKINGS

We use similar commands as used for education center.

OUTPUT (SOME PART ONLY)

186 <http://www.foxsports.com/college-football/>
 113 <http://www.foxsports.com/soccer/>
 97 <http://www.foxsports.com/mlb/>
 93 <http://www.foxsports.com/nfl/>
 93 <http://www.foxsports.com/golf/>
 66 <http://www.isaa.org/popups/vendors/>
 65 <http://www.foxsports.com/nascar/>
 62 <http://www.nblcanada.com/nblcanada/images/>
 45 <http://www.foxsports.com/nba/>
 39 <http://www.foxsports.com/golf/story/>
 36 <http://www.foxsports.com/soccer/story/>
 34 <http://www.foxsports.com/college-football/story/>
 31 <http://www.foxsports.com/nascar/photos/>
 30 <http://www.foxsports.com/>
 29 <http://www.isaa.org/popups/>
 26 <http://www.nbcsports.com/>
 24 <http://www.isaa.org/>
 23 <http://www.nbcsports.com/video/>
 22 <http://www.nfl.com/scores/2015/>
 22 <http://www.foxsports.com/college-football/photos/>
 20 <http://www.nfl.com/>
 20 <http://www.foxsports.com/soccer/lists/>
 19 <http://www.realmadrid.com/>
 19 <http://www.nbcsports.com/node/>
 19 <http://www.fiba.com/File/>
 18 <http://www.ussoccer.com/>
 18 <http://www.foxsports.com/nfl/story/>
 16 <http://www.foxsports.com/mlb/story/>
 14 <http://www.foxsports.com/tag/>
 14 <http://www.foxsports.com/nascar/shake-and-bake/>
 14 <http://www.foxsports.com/motor/story/>
 12 <http://www.foxsports.com/soccer/paper-chase/>

12 <http://www.foxsports.com/nfl/photos/>
12 <http://www.foxsports.com/nba/story/>
12 <http://www.foxsports.com/nascar/story/>
12 <http://www.fiba.com/>
11 <http://www.ussoccer.com/womens-national-team/>
11 <http://www.foxsports.com/college-football/outkick-the-coverage/>
10 <http://www.ussoccer.com/mens-national-team/>
8 <http://www.ussoccer.com/us-under19-mens-national-team/>
8 <http://www.realmadrid.com/zh/football/>
8 <http://www.realmadrid.com/zh/>
8 <http://www.realmadrid.com/ja/football/>
8 <http://www.realmadrid.com/id/>
8 <http://www.realmadrid.com/fr/>
8 <http://www.realmadrid.com/en/football/>
8 <http://www.realmadrid.com/en/>
8 <http://www.realmadrid.com/ar/football/>
8 <http://www.realmadrid.com/ar/>
8 <http://www.nbcsports.com/video/league/>
8 <http://www.fifa.com/ballon-dor/organisation/>
7 http://www.wimbledon.com/en_GB/atoz/
7 <http://www.ussoccer.com/us-under23-mens-national-team/>
7 <http://www.ussoccer.com/us-under20-womens-national-team/>
7 <http://www.ussoccer.com/us-under20-mens-national-team/>
7 <http://www.ussoccer.com/us-under19-womens-national-team/>
7 <http://www.ussoccer.com/us-under18-mens-national-team/>
7 <http://www.realmadrid.com/sobre-el-real-madrid/historia/>
7 <http://www.realmadrid.com/pt/sobre-o-real-madrid/historia/>
7 <http://www.realmadrid.com/pt/futebol/jogos/>
7 <http://www.realmadrid.com/pt/>
7 <http://www.realmadrid.com/ja/>
7 <http://www.realmadrid.com/id/sepak-bola/pertandingan/>
7 <http://www.realmadrid.com/futbol/partidos/>
7 <http://www.realmadrid.com/fr/football/matches/>
7 <http://www.realmadrid.com/fr/a-propos-du-real-madrid/histoire/>
7 <http://www.nbcsports.com/live-extra/>
7 <http://www.isaa.org/pdfs/>
7 <http://www.foxsports.com/nfl/video/>
7 <http://www.foxsports.com/nascar/video/>
7 <http://www.foxsports.com/golf/usga/story/>
7 <http://www.fifa.com/>
6 http://www.wimbledon.com/en_GB/wimbledonfoundation/
6 <http://www.ussoccer.com/us-under23-womens-national-team/>
6 <http://www.ussoccer.com/us-under18-womens-national-team/>
6 <http://www.realmadrid.com/zh/about-real-madrid/history/>
6 <http://www.realmadrid.com/ja/about-real-madrid/history/>
6 <http://www.realmadrid.com/id/basket/pertandingan/>
6 <http://www.realmadrid.com/fr/zone-vip/>
6 <http://www.realmadrid.com/en/vip-area/>
6 <http://www.realmadrid.com/en/basketball/>
6 <http://www.realmadrid.com/en/about-real-madrid/history/>
6 <http://www.realmadrid.com/ar/about-real-madrid/history/>
6 <http://www.foxsports.com/golf/usga/>
6 <http://www.foxsports.com/college-football/lists/>
6 <http://www.fifa.com/womens-football/>
5 <http://www.realmadrid.com/zh/vip-area/>
5 <http://www.realmadrid.com/zh/basketball/games/>
5 <http://www.realmadrid.com/zh/basketball/>
5 <http://www.realmadrid.com/sobre-el-real-madrid/el-club/>
5 <http://www.realmadrid.com/pt/futebol/>
5 <http://www.realmadrid.com/pt/basquetebol/jogos/>
5 <http://www.realmadrid.com/pt/basquetebol/>
5 <http://www.realmadrid.com/pt/area-vip/>
5 <http://www.realmadrid.com/ja/vip-area/>
5 <http://www.realmadrid.com/ja/basketball/games/>
5 <http://www.realmadrid.com/ja/basketball/>
5 <http://www.realmadrid.com/id/tentang-real-madrid/sejarah/>
5 <http://www.realmadrid.com/id/sepak-bola/>
5 <http://www.realmadrid.com/id/keramahtamahan/>
5 <http://www.realmadrid.com/futbol/>

Phase 2 - Search Engine and Indexer

The purpose of second phase was to basically build a search engine and indexer which accepted the input of Phase 1 and provided an Inverted Index and searching and ranking features. Since nutch fetches webpages, therefore the motive behind building an inverted index was to enable search.

An Inverted Index is a key-value index where the key is a word and the value is the document (number) which contains that word, in a corpus (collection of documents). This enables searching - where a user can provide a search term consisting of one or more words which are then fed as input to the search engine and by utilizing the inverted index, the corresponding documents (numbers) are retrieved.

However, the task of the search engine doesn't end here, this is because a word (e.g. the) could potentially exist in too many documents. Therefore, providing an appropriate ordering by using a Ranking Algorithm is required. The objective of the Ranking Algorithm is that, once using the inverted index the documents (numbers) which contain a word have been identified, it ranks the documents (URLs associated with documents) from the most relevant to the least relevant.

Even though the relevance of a document given a search term is a topic, which is beyond the scope of this document, however, it is greatly significant. Without a good ranking algorithm a search engine is, at best, a giant index which gives the user all the documents which contain word(s) entered by the user. One of the key factors which affects the success of a search engine is the effectiveness with which its ranking algorithm can capture the relevance of a document given a search term and rank (order) the documents (URLs associated with the documents).

2.1 Development Environment of the Search Engine and Indexer

The Development Environment for the Search Engine and Indexer consisted of -

1. Programming Language - *Java* SDK 8 (the binaries were compiled to be compatible with JDK 6 because of the JDK 6 available in the Deployment Environment). Java was considered as a natural choice because of the *Write-Once-Run-Anywhere* nature of the Java platform and was further supported by the fact the crawler - nutch is also written in Java (considering future possibility of integrating the two).
2. Integrated Development Environment used - *Eclipse Mars*.

2.2 File Processing

The input from Phase 1 i.e. dump.txt had more than 1.6 million lines of text consisting of all the data, parsed data, urls and metadata and much more information fetched by the crawler. The objective of File Processing was to extract the parsed text of the web pages and the corresponding web pages. So that once we had each web page (with its URL) and the parsed text associated with it, we could start building an inverted-index on it.

It was a difficult task to process files, because there was lot of noise in the form of URLs, metadata, timestamps and information which did not contribute to our goal of extracting web pages' parsed texts and respective URLs.

The dump.txt was first visually inspected and based on visual inspection certain patterns could be observed. The next step was to programmatically implement what had been the observed during visual inspections to extract the parsed web pages' contents and associated URLs

In File Processing -

1. [FileMetadataGenerator.java](#) - accepts any file (obtained from nutch's readseg command i.e. the dump.txt) and generates a metadata which is used in the file processing algorithm.
2. [FileCleaner.java](#) - It accepts the metadata generated by [FileMetadataGenerator.java](#) and the original file and cleans it as per algorithm.
3. [FileSplitter](#) - This file accepts the outputs after *two* passes of metadata generation and cleaning and then splits the result to generate the documents - each of which corresponds to **one web page and has all that page's fetched parsed text**. All the output files are text files.

The output obtained were 687 files. Each of the files corresponds to a specific web-page. A file consists of the URL (of the webpage) and the parsed text contained in the web page.

Note - all the files in the existing implementation were obtained from the crawling and fetching performed in the month of October 2015. The dump.txt used as input belongs to the October 2015. Since then, the content of several web pages has changed. Therefore this needs to be factored in, when looking at results.

Despite using an effective algorithm for file processing, the visual inspection of all 687 parsed text files revealed that 6 had either no data in it or junk values - this makes the algorithm more

than 99% accurate in separating useful data from junk. For the rest the data was perfect as expected.

However, when the above 6 links were visited either they no longer existed or were file links. One of the links had the same data, which we thought of as “junk data” during visual inspection, hence this couldn’t be attributed as deficiency of the algorithm. Rather this should be viewed as proof of effectiveness of the algorithm. If the web page contains junk, which is then parsed by the nutch, the indexer’s job is to merely index the page, it is not designed for linguistic filtering or processing, which is beyond the scope of this course.

Time Consumed - It takes around 1500 ms

OUTPUT - 1 dump.txt (76.2MB) —-> 687 text documents (4.2MB) (at /<directory-root>/files/)

2.3 Token Generation

In lexical analysis, tokenization is the process of breaking a stream of text, up into words, phrases, symbols, or other meaningful elements called tokens. The list of tokens becomes input for further processing such as parsing or text mining.

Since the objective is to create an Inverted Index, we therefore tokenize to efficiently create an index. The purpose of the Inverted Index is to provide a mapping between a word (key) and the document (value(s) - the document number of the document which contains the word) from the corpus (the collection of documents). Tokenization helps in efficient indexing.

For this stage we take the files obtained from the previous stage (File Processing) as input and then for each file -

1. Read the file
2. Create a Set of words in the file. This is because a Set has the mathematical property of having only unique elements, i.e. a Set can not contain an element more than once. A set can not have redundancy by virtue of its mathematical properties. This property of not having redundancy is utilized for fast index creation.
3. Create a new file containing tokens of the input file. These tokens are then used as keys of the Inverted Index.

The reader by think -

Why bother tokenizing ?

Why not use the files directly to form the Inverted-Index ?

Isn't tokenizing an overhead ?

Tokenizing helps to efficiently index because during indexing the token files can be used and the token files efficiently answer a fundamental question -

Que - Does this particular document contain a word (search-key) ?

Ans - Yes/No ?

Which is all that is required to create an Inverted Index. The fact that the tokenized files - files containing only the tokens for a corresponding file obtained from the File Processing stage, do not have redundancy make Indexing faster and efficient. The tokenized files also serve as inputs to the next stage where we describe Stemming.

However, a limitation of the tokenized files is that they can just answer the question about the existence of a word in a document; they *can not reveal the number of occurrences of a word inside a document*. Since the number of occurrences of a word in a document, is an input in many ranking algorithms (e.g. TFIDF , and Document Quality - both of which are used in this search engine), the tokenized files are of no use for such algorithms.

Note - In the first version of the Search Engine - the tokenized files were used during Indexing. However, in the existing version they just act as input for Stemming (described in the next section). In the current version Indexing is combined with tokenization, this helped to reduce the Indexing time from 90s to 10-15s. (6X to 9X times faster).

OUTPUT - Around 600 ms 687 tokenized files 2.9 MB at <directory-root>/tokenj/

2.4 Stemming

Stemming tries to reduce a word to its roots. The purpose is that, it helps to broaden the scope of search in a situation where the word(s), provided as input by the user, are not “as it is” present in the corpus. e.g. biological and biology both have the same lexical root. Thus it could potentially help to locate a word , which has the same root as the word(s) the user is looking for, but is not present “as it is” in the corpus.

`Stemmer.java` is the class responsible for this. This class was obtained from <http://tartarus.org/martin/PorterStemmer/java.txt>

The `Stemmer.java` was appropriately modified to fit our needs. The methods - `stemTokens()` and `createDifferentStemmedTokens()` inside the `SearchEngine.java` class are responsible for stemming. For each tokenFile that is there a stemmedtokenFile is produced. While producing a stemmedtokenFile it is ensured that, if the stemmed token is same as the original token, then it is *not* copied to the stemmedtokenFile. The reason is that, copying the same word only leads to *redundancy* - which *degrades performance* during search through files. During the search when the index does not contain the word being looked for, then these stemmedtokenfiles are used.

Since the stemmed file will contain only those stemmed words which are different than their original non-stemmed versions, some stemmedtokenFiles can be empty. This implies that all words of the original source tokenFile after stemming remained the same and hence none of the stemmed words were copied to the stemmedtokenfile.

The reader may think -

What is the point having an empty stemmedtokenfile ?

Should not such files be discarded ?

The point is that it is not just a file - it acts a metadata. When we look at a parsed web document in the e.g. `<directory-root>/filesj/file.txt` the application guarantees that there would be a corresponding tokenFile.txt and stemmedtokenFile.txt - these two act as metadata for the original parsed web document. This helps during searches and index creation. This could have future applications in further research.

OUTPUT approximately processed in 300 ms at `<directory-root>/differenstemmedtokenj/`

Note- the directory is named as “differentstemmedtoken” - which conveys that in this directory only those tokens are there which after stemming differ from the original token given as input to the process of stemming , as already explained.

2.5 Index Generation

The Index Generation is one of the most critical objectives to be fulfilled. We create an Inverted-Index which is a key-value mapping from words (keys) to document numbers (value). The doc-

ument number would be then used to find the document, in the corpus, and the document would contain the URL, to retrieve which is the ultimate goal of the Indexing.

First the `generateIndices()` method of the `IndexGenerator` class is called. This method calls the `getSortedWordCounter()` method of the same class and the `getSortedWordCounter()` method scans through all parsed file in the `<directory-root>/filesj/` - in this case 687 files to generate a `HASHMAP<STRING,INTEGER>` which maps every distinct string to its number of occurrences across all parsed 687 pages. Then the application sorts them from Least Word Count to High Word Count - ascending order by word count - this hashmap - `sortedWordCounter` is used to calculate - `lowFrequencyWords` and `highFrequencyWords` and `total words`.

During Indexing we maintain two sets `lowFrequencyWords` and `highFrequencyWords`. These are used for Document Quality algorithm that we use during ranking.

1. `lowFrequencyWords` - These are words which occur just 1% in the entire corpus.
2. `highFrequencyWords` - The remaining words which constitute the remaining 99% of the corpus.

Filtration - This increases the efficiency of the indexing . For this, a primitive and augmented filter is used in the same class. The primitive filter consists of digits and alphabets - it is then transformed into an augmented filter by adding all 255 ascii values to it and numbers. `Locations are - Primitive Filter in LINE 29 - IndexGenerator.java and Augmented Filter at LINE 130 and LINE 133 - IndexGenerator.java`. If filtration is not used, then garbage values occupy the index - these values *do not* contribute to the search process.

Statistics about words

1. **Total** number of words in the `sortedWordCounter` = **411598**
2. **High** Frequency words **10780**
3. **Low** Frequency words **6584**
4. Number of **distinct** words = **17364** e.g. the word `the` occurs more than 10000 times

Finally by utilizing the methods of the class `IndexGenerator` the Inverted-Index is generated and is called "primaryIndex.txt" and is stored at `<directory-root>/primaryindexj/`

STATISTICS about primaryindex.txt

1. *TYPE - KEY,VALUE MAPPINGS* key - the word , value - the document number where the word occurs
2. *SIZE - 1.1 MB (NOTE the size of all 687 parsed web pages combined was 4.2 MB) => InvertedIndex here occupies a quarter of the space as all the parsed data.*

3. *LINES* - 34728 which is correct because it is 2 times the number of distinct words in the parsed files - which is 17364. One line to store the word as KEY and another line to store all the document numbers where the word (key) occurs as VALUES.

It is vital for the normal operations of the application that, every time the application runs , it first checks if this file is in its correct location - <directory-root>/primaryindexj/primaryindex.txt If the application does not find the file at its designated location, then it is assumed by the application that it is the first time the application is being run and that indexing never ever took place. Therefore all processes- FILE PROCESSING, TOKEN GENERATION, STEMMING, INDEX GENERATION , SCORING are performed from scratch. If in the beginning , the application detects this file is present at its designated location - it assumes that this file was created during some previous run and the application directly loads the content of the Inverted Index into the main memory for usage during answering search queries.

CONCLUSION - No primaryIndex.txt =>

File Processing, Token Generation,
Stemming, Index Generation,
Scoring done from scratch.

In the First Version of the Search Engine and Indexer - Indexing used to take 10 minutes

In the Second Version

- Indexing used to take about 90 seconds

In the Current Version

- Indexing takes just 10 - 15 seconds

We have been able to speed up the creation of Inverted-Index by about 40 X to 60 X over the previous version.

2.6 Scoring

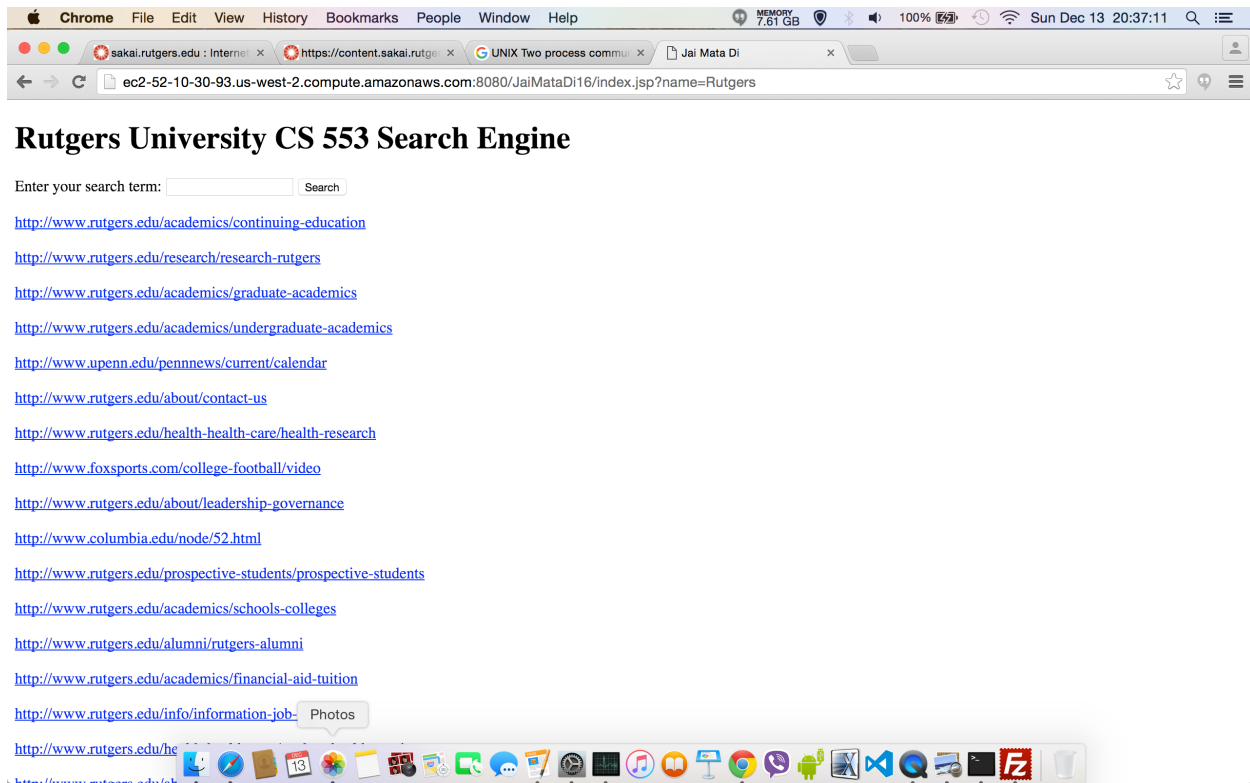
The Scoring is responsible for pre-processing of the documents in corpus. We use a score - Document Quality during ranking. The notion of Document Quality is -

Words which occur less frequently convey more semantic weight (*more notions packed inside a single word*) than frequently occurring words such as “of” , “for” , “the” etc. Therefore we say that words which constitute close to 1% of the total number of words in the entire corpus are **lowFrequencyWords** (See 2.5 for details) and the rest are called **highFrequencyWords**.

Document Quality of a document = **lowFrequencyWords** / (total number of word in a document)

The higher the ratio (on the R.H.S) the higher the Document Quality. DQ constitutes 60% of weightage during assigning total score to a page/document. Since the DQ of a document is independent of a search term, therefore it remains unchanged throughout the lifetime of a document in the corpus. Therefore, we pre-process DQ for all the parsed web pages stored in the *filesj* directory when the application runs for the first time. In subsequent usages these scores are directly loaded in to the main memory. This saves precious CPU cycles. The reader would have noticed that, this technique is similar to the one that is used for Inverted Index, where once the index is created, during the first run of the application, the Inverted Index is thereafter loaded directly into the main memory during subsequent usages.

Phase 3 - Web Interface



The Web-Interface (<http://ec2-52-10-30-93.us-west-2.compute.amazonaws.com:8080/JaiMataDi16/>) was developed using JSP. The usage of the Web-Interface is as follows -

1. The user enters the search term in the text box adjacent to the “Search” button as seen in the screenshot.
2. The search engine generates the output - the urls ranked using the algorithm explained later and displays them on the same page.
3. The flow of search will be explained later.

We share here the information about the Deployment Environment -

1. *Operating System* - Linux
2. *Location* - Cloud - Amazon AWS
3. *Java Environment* - Open JDK 6.0 (On the Server) - The application was developed using JDK 8 and binary was compiled to JDK 6
4. *Web Server* - *GlassFish 3.0 server* (during development GlassFish 4.1 and Netbeans IDE used)

4. Flow and Algorithm

We explain the flow and the algorithm together -

1. The user enters the search term - which can consist of one word or more than one word.
2. The textbox of the front-end accepts the search term and passes it via the JSP web application to the search engine
3. The search term is passed via the following java instruction in the environment-

```
p = Runtime.getRuntime().exec("java -cp /home/ec2-user/project2/SearchEngine/bin  
searchEngineP.SearchEngine " + <search-term>);
```

The above command calls the Environment (in this case the Linux operating system) and causes the operating system to execute the the [SearchEngine](#).

4. The [<search-term>](#) contains the search-term entered by the user which it is passes as arguments to the [SearchEngine](#)
5. The [main\(\)](#) method of the [SearchEngine](#) passes it to the [externalSearch\(\)](#) method which is responsible for accepting input finding documents that contain the word or words of the search term using classes and various methods responsible for the inverted index, stemmed words and for ranking the output.
6. The algorithm inputs the words from the search-term and tries to find a [minimumList](#). For search term with multiple words, [minimumList](#) contains those documents which contain *all the words in the search-term*. For single word search-term this is the list of all documents containing the single word being searched.
7. The algorithm handles the case where [minimumList](#) is empty - which means for multiple-word search-term that there is no single document in corpus which has all the words of the term.
8. It the collects a list of documents which has words (not necessarily all words in term) from the search-term. The ranking algorithm which is explained later will take care to try to rank such documents higher in the output which have most of the words of the search-term. If search-term has [n](#) words and [all n do not occur](#) in a document (which means that [minimumList](#) is empty), then the ranking algorithm will try to give higher preference to such documents which have [n-1 words](#) over documents with [n-2 words from all the n words which constitute the search term](#).
9. *Intuition behind this - If it possible to find all the words which constitute a search term, in one document then rank such a document high in the output list , otherwise if it is not possible to find all words in one document , then the algorithm tries to rank those documents which have most words from the search-term higher in the output, over those documents which have fewer of the words of the search-term. For single word search-term the algorithm either*

finds all those documents in the corpus which contain that word. Thereafter the ranking is done in the following manner using the formula - .

Final Score of Document = 40% TFIDF + 60% DOCUMENT QUALITY (PRE-PROCESSED QUALITY SCORE)

Note - For multiple-words-search-term if there is no document in the corpus, which contains all the words in the search term, even though the algorithm tries to rank a document which has “most” words of the out of all the words (which constitute the multiple word search term), despite this fact it can not guarantee that a document with n-1 words (out of say n words of the multiple-word-search-term) would be ranked higher than a document with n-2 words. This is because 60% of the total score depends on the quality of the document, which is captured by the Document Quality score.

10. TF - IDF calculated using the `Ranker` class `tfidfValue(String keyword, String fileName)` method. The TF-IDF used follows the formula

tfidf of a word ,w, is

$$tf(w) * idf(w)$$

*tf(w) - (number of times a word appears in a document) / (total number of words in document)
idf(w) - log(Number of documents/number of documents that contain word w)*

11. Document Quality - Calculated using `score(int docNumber)` method of the `Ranker` class.
12. During Indexing we maintain two sets `lowFrequencyWords` and `highFrequencyWords`. They are stored in files for future usage.
13. `lowFrequencyWords` - These are words which occur just 1% in the entire corpus.
`highFrequencyWords` - The remaining words which constitute the remaining 99% of the corpus.
14. After indexing is over, the scoring of all documents in the corpus is done as a part of pre-processing, and the scores are stored and are loaded to main memory in subsequent usage.

Just like indexing is done once and stored so that during subsequent usage it can be loaded to main memory, these scores are also stored and later loaded into the main memory using a data structure - hash table which maintains the scores of all the documents of the corpus.

The pre-processing of scores is done to save time and resources during subsequent usage. As we will explain below, the quality of document once fetched using crawler remains unchanged throughout its lifetime in the search engine, because it is the quality of the document as a whole and is completely independent of the search term.

The scoring is an expensive operation, since it is done once, when the application runs for the first time and is subsequently stored, the one-time cost is amortized over subsequent uses.

15. Intuition behind scoring (DQ) - words which occur less frequently in the corpus convey more semantic value than words which occur more frequently in the corpus.
16. A document with higher ratio of `lowFrequencyWords / (total words in that document)` conveys more semantic value in fewer words and therefore is considered as a **higher quality document** than a document which conveys same semantic value in more words i.e. which contains same number of `lowFrequencyWords` however, has higher total number of words. Therefore if the `lowFrequency/Total word count` ratio is higher then the quality of the document is expected to be higher.
17. Thus both `tfidf()` and **Document Quality** is considered while scoring and determining final ranks of the URLs which would be displayed in the output to the user.
18. Then the documents are ranked from highest to lowest based on the score obtained from **40 % weight TFIDF + 60 % Document Quality**
19. *The values 40% and 60% are obtained by empirical observation.*
20. Output is then displayed on the screen.

5. General Observations

1. When the user entered search term containing “,” the comma is converted to %2C in the browser and this degraded the performance and generated poor results.
2. Hence when the user enters comma - “,” or full stop “.” or hyphen “-“. These are removed from the search-term by the Web Interface before passing the arguments to the search engine.
3. *Since the previous versions the Indexing performance is better than earlier indexing, because the latter used to consume 10 minutes. After improvements, it used to take 90s and currently it takes just 10-15s to index all the parsed documents in the corpus. **The current indexing time is 40X to 60X times faster than its performance during first version.***
4. The Document Quality improved ranking over just using TFIDF alone to determine the final score of a document.
5. We believe that since Web-Interface calls the process - search engine through the environment this could be more amenable for cloud implementation of this project because in the cloud computing model, a organization is charged by cloud's service provider, by the number of processors and memory used. As many SearchEngine processes will be launched as the number of users, which means the *usage of memory would grow linearly with number of users.*
6. The indices of the SearchEngine use HashMaps of Java and are efficiently algorithmically designed to be minimal in size (the size of index when on disk is just 1.1 MB for 687 web pages). This implies that for a search engine which has data of about 100,000 high quality web pages, the size of the index would be about 155 MB (assuming similar distribution of words as the given corpus of parsed documents.)

Note - Currently even the commodity hardware computers have 8 GB RAM , therefore cloud computing is hoped to have better resources - on pay-as-you-use basis. Therefore this architecture should be scalable and would not be constrained by sockets.

7. It was found that nutch crawled several thousand links (see above for more details) but the output of nutch does not include parsed data for all the links crawled by it. When the dump.txt was visually inspected there were instances of pages that had certain fraction of the page's data missing. This could be attributed to packet loss during transmission while fetching
8. The delay introduced in crawling is mainly due to the *robot.txt* files set by system administrators of the websites. This leads to slow crawling. It could be observed that during certain instances of crawling, the speed of fetching documents from a website could be as slow as *6 pages / minute*. This implies that, to crawl a web site with 100 web pages at the mentioned rate, it could take up to *17 minutes*. Considering there are thousands of websites, this is speed is slow.

9. *Technical Challenges* - We need to provide the path to the search engine. Initially, it was assumed that the search engine would be able to automatically find out its location in the deployment environment using `getAbsolutePath()` or `getCanonicalPath()` methods of `File` class of Java. However, this has not been successful in the deployment environment, though in development environment this worked as expected. *We think, that this should be due to the in-built security features of the Operating Systems offered by Amazon's cloud. This is because if a file (written in java or any other language) can automatically find its own location, then such an information puts it in a significant position to cause damage, if the file happens to be written for malicious purposes.* It is envisioned to overcome this challenge, by utilizing a separate configuration file for the search engine, which would be utilized to read the path to the search engine (as is the case with several commercial software).

Summary of Contributions

Phase 1 - Mrinal Monga - Understanding the nuances of nutch. Crawling and analyzing crawled data - which is used currently by search engine.

Riham S - Crawled and text processing

Phase 2 - Mrinal Monga - Designing and Developing the Search Engine and Indexer in Java.

Phase 3 - Mrinal Monga - Web-Interfacing - Designed and Developed the existing UI and interfacing used between UI and the Search Engine and Indexer in the Deployment Environment. Deployment of the Search Engine and Indexer and UI and interfacing in the Deployment Environment.

Riham S - Worked on a different UI.